

The Central Limit Theorem - simulation of random exponential samples

Synopsis

This project is about simulation of random variables and the central limit theorem. The means of 1000 samples of 40 iid random variables (drawn from an exponential distribution) are treated as a random sample, and compared to a normal distribution. ## Exponential distribution

For random variables drawn from a (continuous) exponential distribution of rate λ , the propability distribution function is given by

$$P(x) = \lambda e^{-\lambda x}, \quad x \geq 0.$$

This distribution has a population mean of $\mu = 1/\lambda$ a population variance of $1/\lambda^2$, and a standard deviation of $\sigma = 1/\lambda$.

Central Limit Theorem

The central limit theorem states that as the sample size increases the mean of a sample of iid random variables should itself behave like a random normal variable, with mean equal to the mean of the original population, and standard deviation equal to the standard error of the sample mean.

Simulation

```
#opts_chunk$set(fig.width=8, fig.height=8, dpi=144)
#https://github.com/lgreski/datasciencecontent/blob/master/markdown/statinf-expDistCheck
require(ggplot2)

## Loading required package: ggplot2

require(dplyr)

## Loading required package: dplyr

##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```

require(gridExtra)

## Loading required package: gridExtra

##
## Attaching package: 'gridExtra'

## The following object is masked from 'package:dplyr':
##
##      combine

generate two datasets, the first is a sample of 1000 iid random variables drawn
from an exponential distribution with rate  $\lambda = 0.2$ . The second dataset consists
of 1000 sample means, where each sample consists of 40 iid random variables
drawn from an exponential distribution with rate  $\lambda = 0.2$ 

set.seed(5074491)
samples<-1000
lambda<-0.2
samples1k<-rexp(samples,rate=lambda)
means1kx40<-replicate(samples,mean(rexp(40,rate=lambda)))
# data1k<- data.frame(samples1k)
# names(data1k)<-'expSamples'
# data1kx40<-data.frame(means1kx40)
# names(data1kx40)<-'meanExp40Samples'

#names(data)<-c('expSamples','meanExp40Samples')

```

sample mean

The population mean for an exponential distribution is equal to $1/\lambda$, which in this case is $1/0.2 = 5$. For each sample, the mean should be close to this value. The mean of 1000 sample means should also be close to this value. For the samples generated for this report, the mean is given by

```

themean<-mean(means1kx40)
themean

## [1] 5.000418

```

sample variance

As the sample under consideration is a set of means, the variance of this sample should be given by

$$var(\bar{X}) = \frac{\sigma^2}{n},$$

which is the square of the standard error in the mean. For the exponential distribution, the variance (σ^2) is given by $1/\lambda^2$. The expected variance for our sample of means would then be

$$\text{var}(\bar{X}) = \frac{1}{\lambda^2 n}$$

For samples of 40 variables, with $\lambda = 0.2$, we would expect the means to have variance 0.625

```
var(means1kx40)
## [1] 0.5938206
```

The value here is 0.594, which is pretty close to the expected 0.625. Another way to check that the distribution is normal is to evaluate the probability of a variable lying within a certain range of the mean. For a normal distribution, 68% of the values should lie within one standard deviation of the mean. For our sample of exponential means, this information is summarised in the table below

X	Area within $X\sigma$ of mean for a normal distribution	fraction within $X\sigma$ for sample of means
1	0.683	0.706
2	0.954	0.957
3	0.997	0.998
4	1.000	1.000

```
sigmas<-c(1,2,3,4)
fractions<-sapply(sigmas, function(x) { sum((means1kx40 < (themean + x/(lambda*sqrt(40))) ))
round(fractions,digits=3)

## [1] 0.706 0.957 0.998 1.000

# normal distribution
normfracs<-sapply(sigmas,function(x) { pnorm(x) - pnorm(-x)})
round(normfracs,digits=3)

## [1] 0.683 0.954 0.997 1.000

# plot distributions
dnormscale<-function(x,scale=1.0,...) { scale*dnorm(x,...)}
dexpscale<-function(x,scale=1.0,...) { scale*dexp(x,...)}
df<-data.frame(samples1k,means1kx40)
hist_exp<-ggplot(data=df,aes(x=samples1k)) + geom_histogram(binwidth=0.2,aes(fill=..count..),
scale_fill_gradientn(colours=c('blue','purple')) + stat_function(fun=dexpscale,colour='red',
hist_expMean<-ggplot(data=df,aes(x=means1kx40)) + geom_histogram(binwidth=0.2,aes(fill=..count..))
```

```
labs(title='distribution of 1000 sample means of 40 random exponential variables',x='value',
stat_function(fun=dnormscale,colour='red',args=list(scale=0.2*samples,mean=1/lambda,sd=1/(1a

#hist_exp
grid.arrange(hist_exp,hist_expMean,nrow=1)
```




Figure 1: plot of chunk plot

add curves to these, the population distribution and also the exponential and gaussian f

Part 1: Simulation Exercise Instructions In this project you will investigate the exponential distribution in R and compare it with the Central Limit Theorem. The exponential distribution can be simulated in R with `rexp(n, lambda)` where `lambda` is the rate parameter. The mean of exponential distribution is $1/\lambda$ and the standard deviation is also $1/\lambda$. Set $\lambda = 0.2$ for all of the simulations. You will investigate the distribution of averages of 40 exponentials. Note that you will need to do a thousand simulations.

Illustrate via simulation and associated explanatory text the **properties of the distribution of the mean of 40 exponentials** . You should

Show the sample mean and compare it to the theoretical mean of the distribution. Show how variable the sample is (via variance) and compare it to the theoretical variance of the distribution. Show that the distribution is approximately normal.